

AAI_510_5_final_Project

June 12, 2025

1 Predicting Term Deposit Subscription Using Machine Learning

1.1 Introduction

In today's data-driven banking environment, improving the effectiveness of customer outreach is both a strategic goal and an operational necessity. Traditional marketing campaigns for term deposit products often involve contacting thousands of customers with only a small fraction converting into actual subscribers. This results in low efficiency, increased costs, and potential customer fatigue.

This project aims to solve that problem by using machine learning to predict which customers are most likely to subscribe to a term deposit before making contact. We leveraged customer demographic data, past campaign outcomes, economic indicators, and behavioral patterns to build a predictive model that enhances marketing precision.

Our solution not only improves campaign return on investment (ROI) but also ensures ethical targeting and supports regulatory compliance. This report outlines our end-to-end process, key findings, and actionable business recommendations in a format tailored for non-technical stakeholders.

```
[5]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
[6]: df = pd.read_csv('bank-additional-full.csv', delimiter = ';')
```

```
[7]: df.head()
```

```
[7]:   age      job marital  education default housing loan  contact \
0   56  housemaid  married   basic.4y      no      no   no  telephone
1   57  services  married high.school unknown      no   no  telephone
2   37  services  married high.school      no   yes   no  telephone
3   40   admin.  married   basic.6y      no      no   no  telephone
4   56  services  married high.school      no      no  yes  telephone

   month day_of_week ... campaign pdays previous      poutcome emp.var.rate \
```

0	may	mon	...	1	999	0	nonexistent	1.1
1	may	mon	...	1	999	0	nonexistent	1.1
2	may	mon	...	1	999	0	nonexistent	1.1
3	may	mon	...	1	999	0	nonexistent	1.1
4	may	mon	...	1	999	0	nonexistent	1.1

	cons.price.idx	cons.conf.idx	euribor3m	nr.employed	y
0	93.994	-36.4	4.857	5191.0	no
1	93.994	-36.4	4.857	5191.0	no
2	93.994	-36.4	4.857	5191.0	no
3	93.994	-36.4	4.857	5191.0	no
4	93.994	-36.4	4.857	5191.0	no

[5 rows x 21 columns]

```
[8]: print("Total Columns", df.shape[1])
```

Total Columns 21

```
[9]: features = df.drop(columns=['y'])
print("Number of features: ", features.shape[1])
```

Number of features: 20

We only want to keep original 16 features

```
[10]: # define original 16 features from df
original_16_features = [
    'age', 'job', 'marital', 'education', 'default', 'housing', 'loan',
    'contact', 'month', 'duration', 'campaign', 'pdays', 'previous',
    'poutcome', 'emp.var.rate', 'cons.price.idx'
]

# create new df with only these features and target variable
df_16 = df[original_16_features + ['y']]
```

```
[11]: df_16.head()
```

	age	job	marital	education	default	housing	loan	contact	\
0	56	housemaid	married	basic.4y	no	no	no	telephone	
1	57	services	married	high.school	unknown	no	no	telephone	
2	37	services	married	high.school	no	yes	no	telephone	
3	40	admin.	married	basic.6y	no	no	no	telephone	
4	56	services	married	high.school	no	no	yes	telephone	

	month	duration	campaign	pdays	previous	poutcome	emp.var.rate	\
0	may	261	1	999	0	nonexistent	1.1	
1	may	149	1	999	0	nonexistent	1.1	
2	may	226	1	999	0	nonexistent	1.1	

3	may	151	1	999	0	nonexistent	1.1
4	may	307	1	999	0	nonexistent	1.1

	cons.price.idx	y
0	93.994	no
1	93.994	no
2	93.994	no
3	93.994	no
4	93.994	no

```
[12]: # seperate features and target
X_16 = df_16.drop(columns=['y'])
y_16 = df_16['y']
```

```
[13]: # check for missing values
print(X_16.isna().sum())
```

```
age          0
job          0
marital      0
education    0
default      0
housing      0
loan         0
contact      0
month        0
duration     0
campaign     0
pdays       0
previous     0
poutcome     0
emp.var.rate 0
cons.price.idx 0
dtype: int64
```

Model 1 Full One Hot Encode brings features to 55. (Better for tree-based, trees handle redundancy well)

Model 2 Reduced One Hot Encode (Better for clustering, less dimensional noise)

```
[14]: # id and one hot encode categorical variables
categorical_cols = X_16.select_dtypes(include= ['object']).columns.tolist()
model1 = pd.get_dummies(X_16, columns=categorical_cols, drop_first=False)
model2 = pd.get_dummies(X_16, columns=categorical_cols, drop_first=True)
```

```
[15]: print(model1.shape[1])
print(model2.shape[1])
```

55

```
[16]: model1.head()
```

```
[16]:   age  duration  campaign  pdays  previous  emp.var.rate  cons.price.idx  \
0    56      261         1    999         0         1.1         93.994
1    57      149         1    999         0         1.1         93.994
2    37      226         1    999         0         1.1         93.994
3    40      151         1    999         0         1.1         93.994
4    56      307         1    999         0         1.1         93.994

   job_admin.  job_blue-collar  job_entrepreneur  ...  month_jul  month_jun  \
0         False              False              False  ...    False    False
1         False              False              False  ...    False    False
2         False              False              False  ...    False    False
3          True              False              False  ...    False    False
4         False              False              False  ...    False    False

   month_mar  month_may  month_nov  month_oct  month_sep  poutcome_failure  \
0         False      True      False      False      False              False
1         False      True      False      False      False              False
2         False      True      False      False      False              False
3         False      True      False      False      False              False
4         False      True      False      False      False              False

   poutcome_nonexistent  poutcome_success
0                   True              False
1                   True              False
2                   True              False
3                   True              False
4                   True              False

[5 rows x 55 columns]
```

```
[17]: model2.head()
```

```
[17]:   age  duration  campaign  pdays  previous  emp.var.rate  cons.price.idx  \
0    56      261         1    999         0         1.1         93.994
1    57      149         1    999         0         1.1         93.994
2    37      226         1    999         0         1.1         93.994
3    40      151         1    999         0         1.1         93.994
4    56      307         1    999         0         1.1         93.994

   job_blue-collar  job_entrepreneur  job_housemaid  ...  month_dec  \
0         False              False              True  ...    False
1         False              False              False  ...    False
2         False              False              False  ...    False
3         False              False              False  ...    False
```

4	False	False	False	...	False	
	month_jul	month_jun	month_mar	month_may	month_nov	month_oct \
0	False	False	False	True	False	False
1	False	False	False	True	False	False
2	False	False	False	True	False	False
3	False	False	False	True	False	False
4	False	False	False	True	False	False

	month_sep	poutcome_nonexistent	poutcome_success
0	False	True	False
1	False	True	False
2	False	True	False
3	False	True	False
4	False	True	False

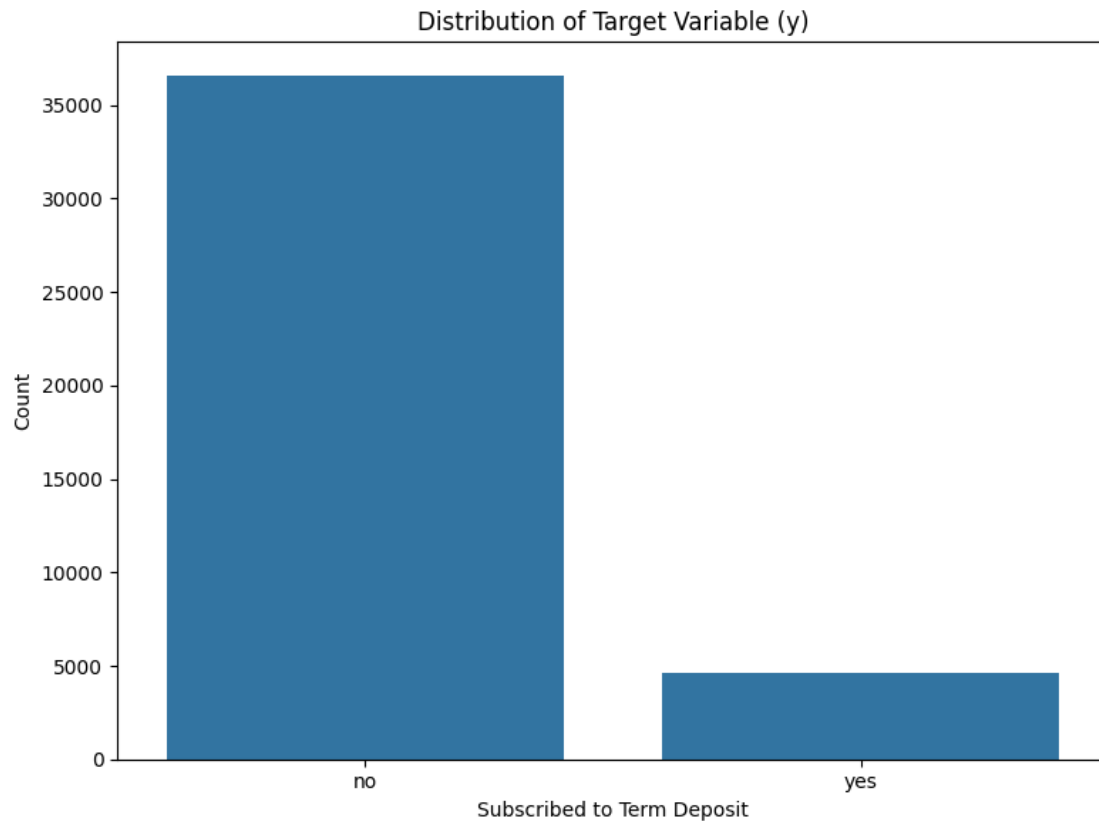
[5 rows x 46 columns]

```
[18]: # Save to csv for model 1 supervised model
df_model1 = model1.copy()
df_model1['y'] = y_16
df_model1.to_csv('df_model1.csv', index=False)
```

```
[19]: # Save to csv for model 2 unsupervised model
df_model2 = model2.copy()
model2.to_csv('df_model2.csv', index=False)
```

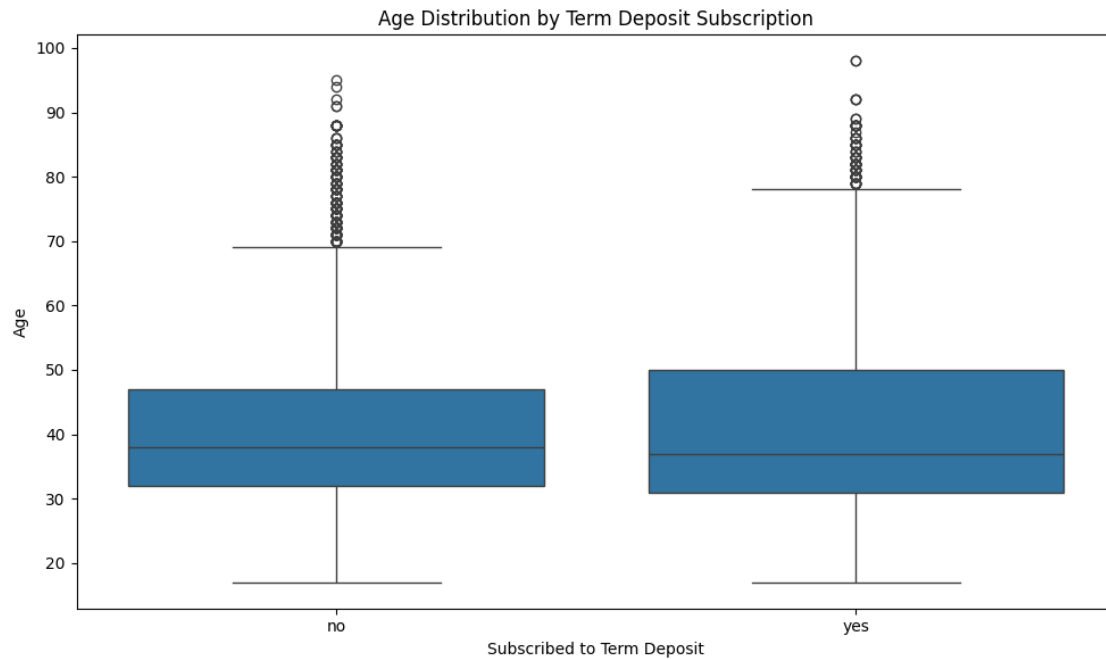
Now let's perform some EDA Distribution of Target Variable (y)

```
[20]: # Distribution of target variable
plt.figure(figsize=(8,6))
sns.countplot(x=y_16)
plt.title("Distribution of Target Variable (y)")
plt.xlabel("Subscribed to Term Deposit")
plt.ylabel("Count")
plt.tight_layout()
plt.show()
```



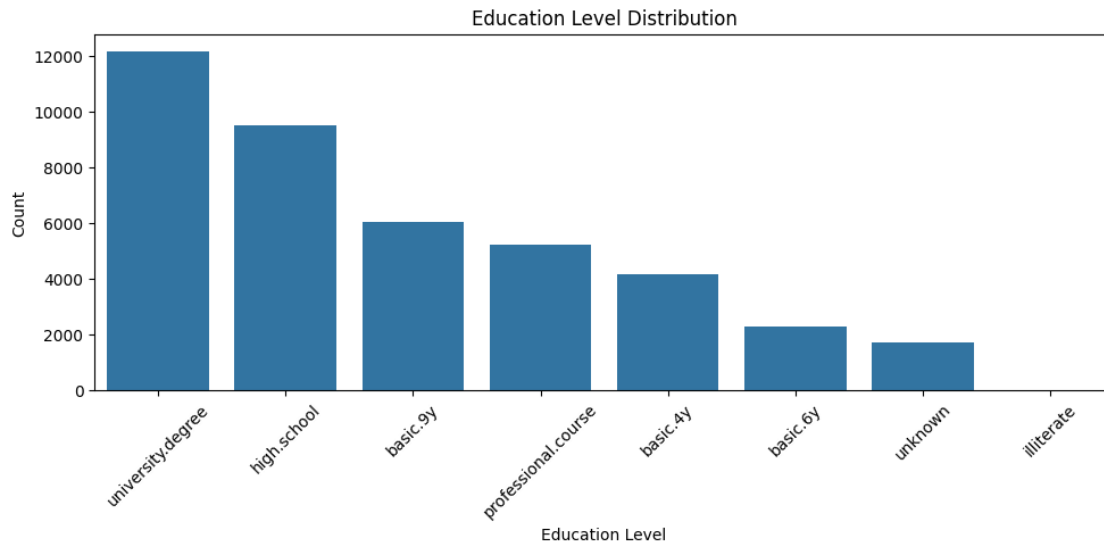
Boxplot of age by Target Variable

```
[21]: plt.figure(figsize=(10,6))
sns.boxplot(x=y_16, y=X_16['age'])
plt.title("Age Distribution by Term Deposit Subscription")
plt.xlabel("Subscribed to Term Deposit")
plt.ylabel("Age")
plt.tight_layout()
plt.show()
```



Countplot of education levels

```
[22]: plt.figure(figsize=(10,5))
sns.countplot(data=X_16, x='education', order=X_16['education'].value_counts().
            ↪index)
plt.title("Education Level Distribution")
plt.xticks(rotation=45)
plt.xlabel("Education Level")
plt.ylabel("Count")
plt.tight_layout()
plt.show()
```

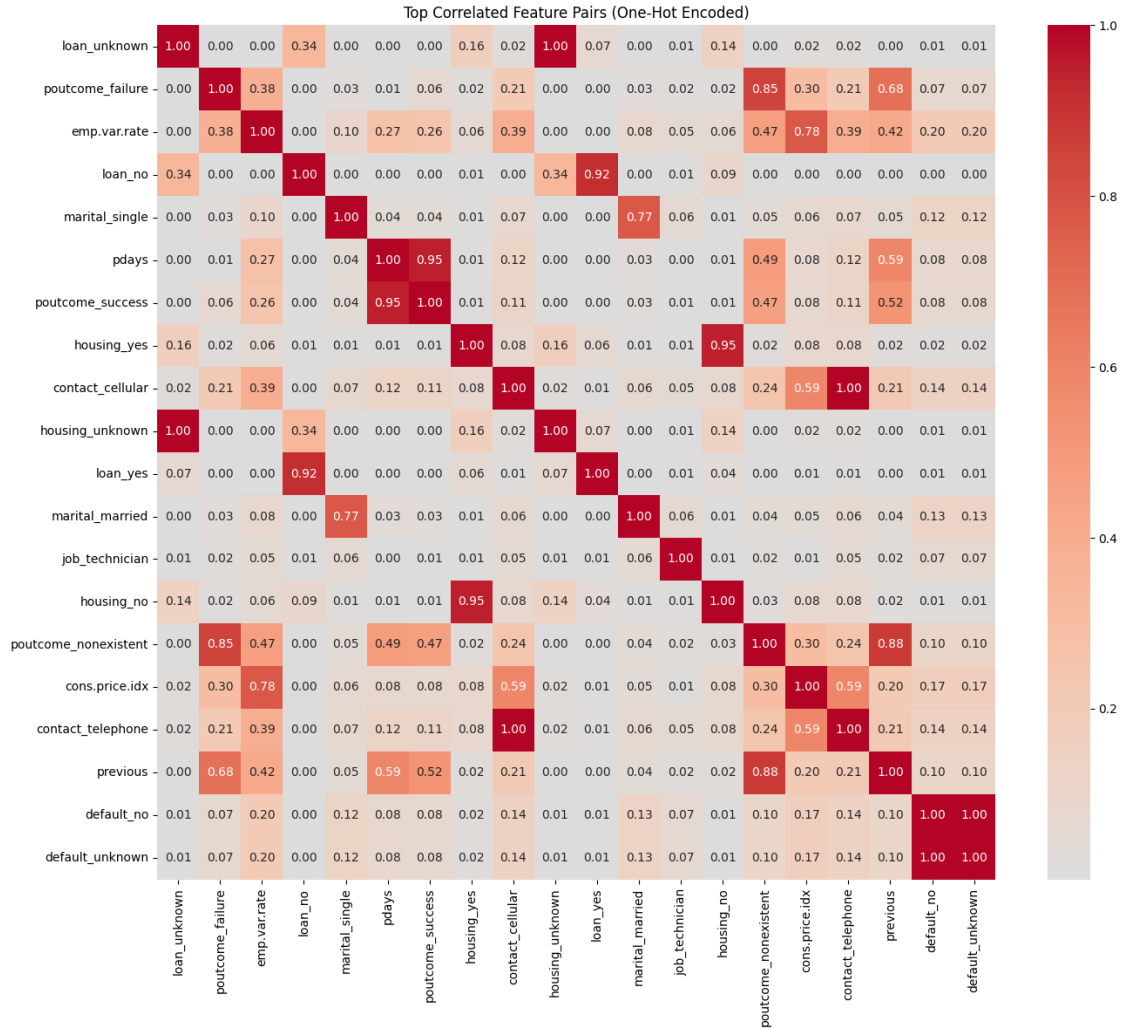


Top 20 Correlation Heatmap of All Encoded Features

```
[24]: plt.figure(figsize=(14, 12))
corr_matrix = model1.corr().abs()
mask = np.triu(np.ones(corr_matrix.shape), k=1).astype(bool)
upper_tri = corr_matrix.where(mask)

sorted_pairs = (
    upper_tri.stack().sort_values(ascending=False)
)
top_features = set()
for i, j in sorted_pairs.index:
    if len(top_features) >= 20:
        break
    top_features.add(i)
    if len(top_features) >= 20:
        break
    top_features.add(j)

filtered_corr = corr_matrix.loc[list(top_features), list(top_features)]
sns.heatmap(filtered_corr, annot=True, cmap="coolwarm", fmt='.2f', center=0)
plt.title("Top Correlated Feature Pairs (One-Hot Encoded)")
plt.tight_layout()
plt.show()
```

```
[25]: df_model1
```

```
[25]:
```

	age	duration	campaign	pdays	previous	emp.var.rate	cons.price.idx	\
0	56	261	1	999	0	1.1	93.994	
1	57	149	1	999	0	1.1	93.994	
2	37	226	1	999	0	1.1	93.994	
3	40	151	1	999	0	1.1	93.994	
4	56	307	1	999	0	1.1	93.994	
...	
41183	73	334	1	999	0	-1.1	94.767	
41184	46	383	1	999	0	-1.1	94.767	
41185	56	189	2	999	0	-1.1	94.767	
41186	44	442	1	999	0	-1.1	94.767	
41187	74	239	3	999	1	-1.1	94.767	

	job_admin.	job_blue-collar	job_entrepreneur	...	month_jun	\
0	False	False	False	...	False	
1	False	False	False	...	False	
2	False	False	False	...	False	
3	True	False	False	...	False	
4	False	False	False	...	False	
...	
41183	False	False	False	...	False	
41184	False	True	False	...	False	
41185	False	False	False	...	False	
41186	False	False	False	...	False	
41187	False	False	False	...	False	

	month_mar	month_may	month_nov	month_oct	month_sep	\
0	False	True	False	False	False	
1	False	True	False	False	False	
2	False	True	False	False	False	
3	False	True	False	False	False	
4	False	True	False	False	False	
...	
41183	False	False	True	False	False	
41184	False	False	True	False	False	
41185	False	False	True	False	False	
41186	False	False	True	False	False	
41187	False	False	True	False	False	

	poutcome_failure	poutcome_nonexistent	poutcome_success	y
0	False	True	False	no
1	False	True	False	no
2	False	True	False	no
3	False	True	False	no
4	False	True	False	no
...
41183	False	True	False	yes
41184	False	True	False	no
41185	False	True	False	no
41186	False	True	False	yes
41187	True	False	False	no

[41188 rows x 56 columns]

```
[26]: df_model2
```

```
[26]:
```

	age	duration	campaign	pdays	previous	emp.var.rate	cons.price.idx	\
0	56	261	1	999	0	1.1	93.994	
1	57	149	1	999	0	1.1	93.994	
2	37	226	1	999	0	1.1	93.994	

3	40	151	1	999	0	1.1	93.994
4	56	307	1	999	0	1.1	93.994
...
41183	73	334	1	999	0	-1.1	94.767
41184	46	383	1	999	0	-1.1	94.767
41185	56	189	2	999	0	-1.1	94.767
41186	44	442	1	999	0	-1.1	94.767
41187	74	239	3	999	1	-1.1	94.767

	job_blue-collar	job_entrepreneur	job_housemaid	...	month_dec	\
0	False	False	True	...	False	
1	False	False	False	...	False	
2	False	False	False	...	False	
3	False	False	False	...	False	
4	False	False	False	...	False	
...	
41183	False	False	False	...	False	
41184	True	False	False	...	False	
41185	False	False	False	...	False	
41186	False	False	False	...	False	
41187	False	False	False	...	False	

	month_jul	month_jun	month_mar	month_may	month_nov	month_oct	\
0	False	False	False	True	False	False	
1	False	False	False	True	False	False	
2	False	False	False	True	False	False	
3	False	False	False	True	False	False	
4	False	False	False	True	False	False	
...	
41183	False	False	False	False	True	False	
41184	False	False	False	False	True	False	
41185	False	False	False	False	True	False	
41186	False	False	False	False	True	False	
41187	False	False	False	False	True	False	

	month_sep	poutcome_nonexistent	poutcome_success
0	False	True	False
1	False	True	False
2	False	True	False
3	False	True	False
4	False	True	False
...
41183	False	True	False
41184	False	True	False
41185	False	True	False
41186	False	True	False
41187	False	False	False

[41188 rows x 46 columns]

```
[27]: print("Total Columns Model 1 df: ", df_model1.shape[1])
      print("Total Columns Model 2 df: ", df_model2.shape[1])
```

Total Columns Model 1 df: 56

Total Columns Model 2 df: 46

```
[31]: df_model_hot_encoded = df_model1
```

```
[32]: df_model_hot_encoded.head()
```

```
[32]:
```

	age	duration	campaign	pdays	previous	emp.var.rate	cons.price.idx	\
0	56	261	1	999	0	1.1	93.994	
1	57	149	1	999	0	1.1	93.994	
2	37	226	1	999	0	1.1	93.994	
3	40	151	1	999	0	1.1	93.994	
4	56	307	1	999	0	1.1	93.994	

	job_admin.	job_blue-collar	job_entrepreneur	...	month_jun	month_mar	\
0	False	False	False	...	False	False	
1	False	False	False	...	False	False	
2	False	False	False	...	False	False	
3	True	False	False	...	False	False	
4	False	False	False	...	False	False	

	month_may	month_nov	month_oct	month_sep	poutcome_failure	\
0	True	False	False	False	False	
1	True	False	False	False	False	
2	True	False	False	False	False	
3	True	False	False	False	False	
4	True	False	False	False	False	

	poutcome_nonexistent	poutcome_success	y
0	True	False	no
1	True	False	no
2	True	False	no
3	True	False	no
4	True	False	no

[5 rows x 56 columns]

```
[36]: cols_to_plot = [
      'age', # age of client
      'duration', # last contact duration in seconds (very skewed, most
      ↪ predictive)
      'campaign', # number of contacts during campaign
```

```

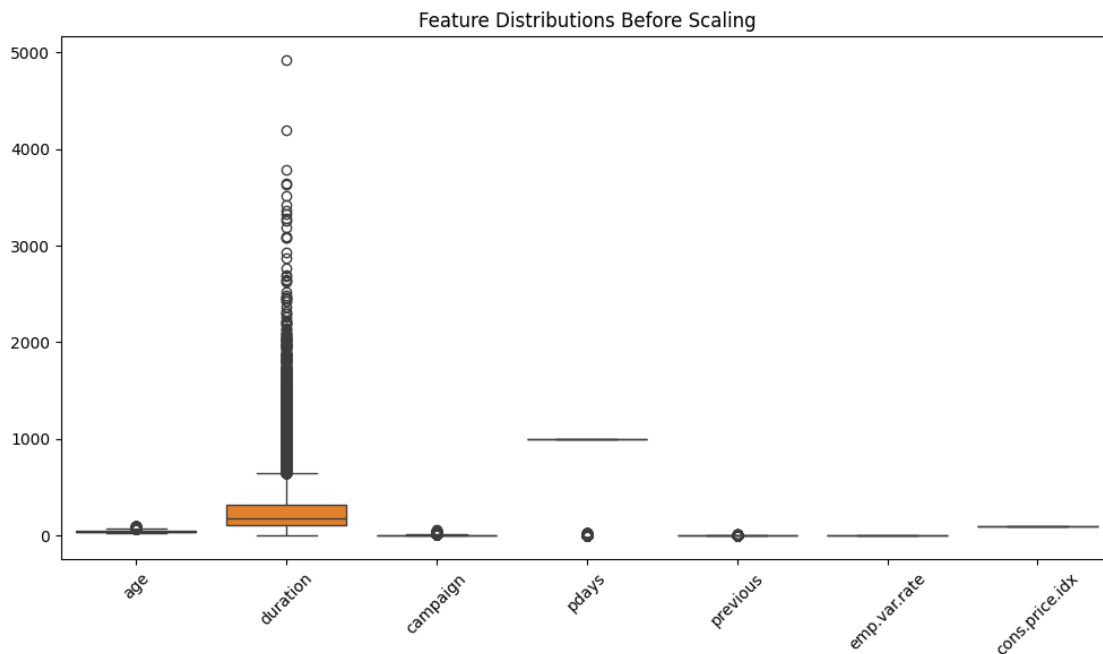
'pdays', # days since last contact (999 = never contacted)
'previous', # number of contacts before this campaign
'emp.var.rate', # employment variation rate (economic indicator)
'cons.price.idx' # consumer price index (economic indicator)
]

```

```

[37]: plt.figure(figsize=(12, 6))
sns.boxplot(data=df_model_hot_encoded[cols_to_plot])
plt.title("Feature Distributions Before Scaling")
plt.xticks(rotation=45)
plt.show()

```

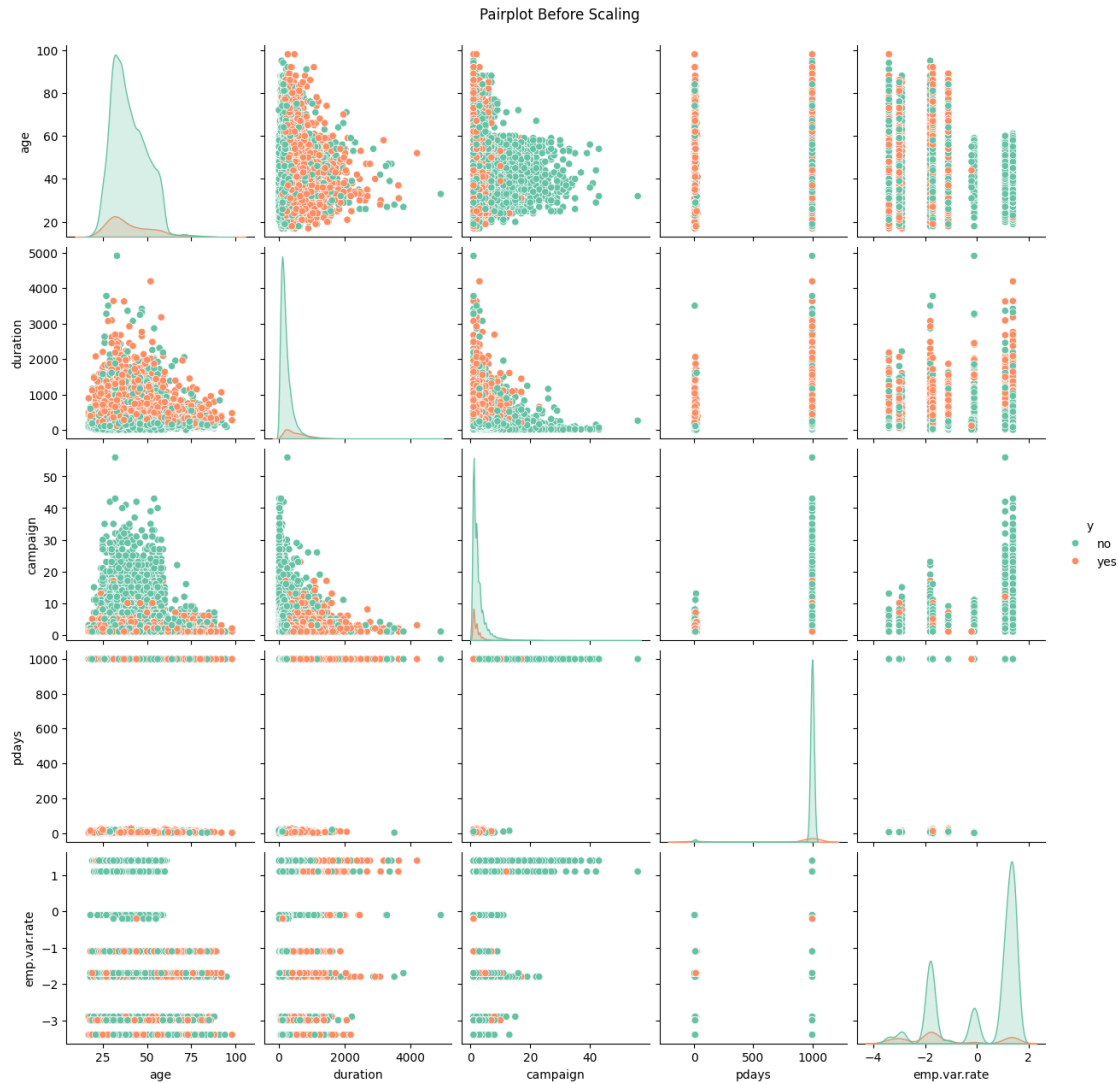


```

[38]: import seaborn as sns
import matplotlib.pyplot as plt
# Subset dataframe
cols_pairplot = ['age', 'duration', 'campaign', 'pdays', 'emp.var.rate', 'y']
df_pair = df_model_hot_encoded[cols_pairplot]

# Plot
sns.pairplot(df_pair, hue='y', palette='Set2', diag_kind='kde')
plt.suptitle("Pairplot Before Scaling", y=1.02)
plt.show()

```



```
[33]: # Separate features and target variable for the hot encoded model
X = df_model_hot_encoded.drop('y', axis=1)
y = df_model_hot_encoded['y']
```

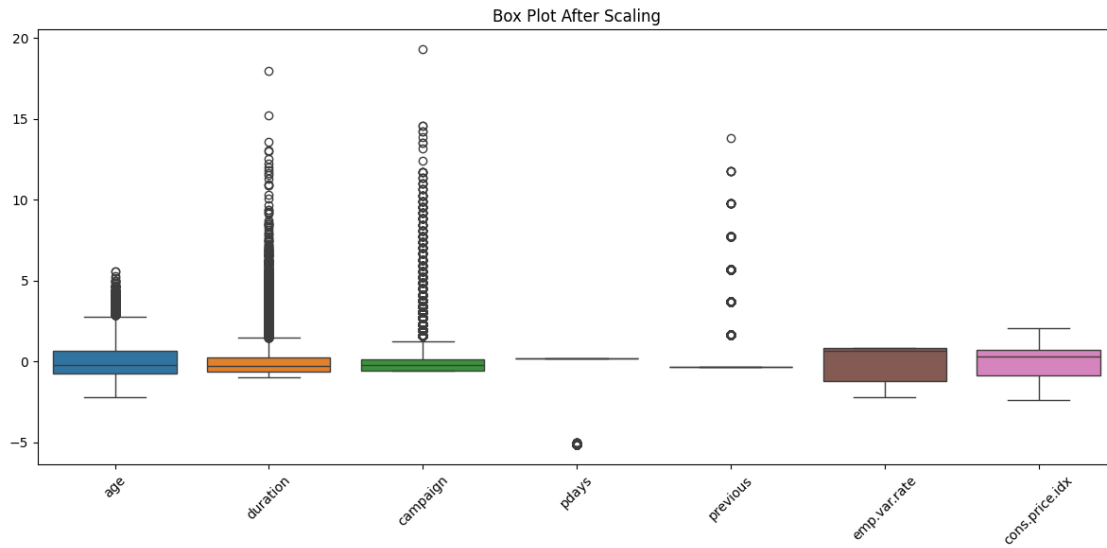
```
[41]: from sklearn.preprocessing import StandardScaler
import pandas as pd

# Scale selected features only
scaler = StandardScaler()
X_scaled_subset = pd.DataFrame(scaler.fit_transform(X[cols_to_plot]),
                               columns=cols_to_plot)

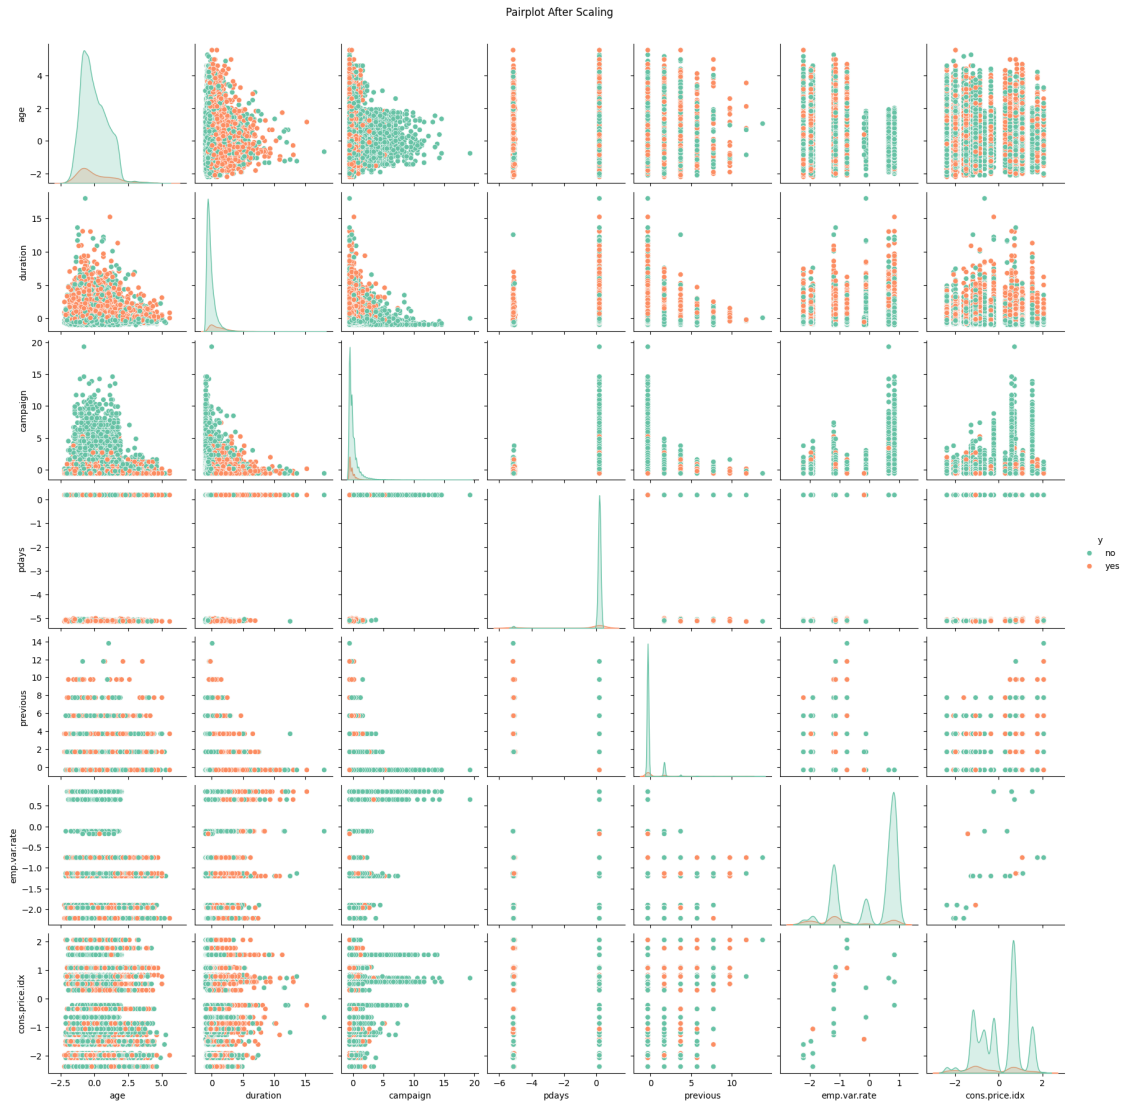
# Add target variable for pairplot
X_scaled_subset['y'] = y.values
```

```
[42]: import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(12, 6))
sns.boxplot(data=X_scaled_subset[cols_to_plot])
plt.title("Box Plot After Scaling")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
[43]: sns.pairplot(X_scaled_subset, hue='y', palette='Set2', diag_kind='kde')
plt.suptitle("Pairplot After Scaling", y=1.02)
plt.show()
```



```
[44]: # Scale the features using StandardScaler
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
[46]: from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
import pandas as pd
import matplotlib.pyplot as plt

# KMeans wrapper function
```



```
def kmeans_execution(df, num_clust):
    kmn = KMeans(n_clusters=num_clust, n_init='auto', random_state=0)
    kmn.fit(df)
    return kmn, kmn.labels_, kmn.inertia_
```

1.1.1 Why We Used Clustering (KMeans)

Before building the prediction model, we used **clustering** to group customers based on similar behaviors and characteristics. This helped us uncover hidden patterns in the data — such as groups of customers who behave alike but respond differently to marketing.

By assigning each customer to a cluster, we created a new feature that adds strategic value to the model. This helps the bank better understand customer segments and tailor outreach strategies accordingly.

```
[47]: # Fill missing values with mean for scaled data
scaled_df_filled = pd.DataFrame(X_scaled).fillna(X_scaled.mean())

k_values = range(2, 7)
silhouette_scores = {}
kmeans_models = {}
kmeans_labels = {}

for k in k_values:
    model, labels, inertia = kmeans_execution(scaled_df_filled, k)
    score = silhouette_score(scaled_df_filled, labels)

    silhouette_scores[k] = score
    kmeans_models[k] = model
    kmeans_labels[k] = labels

    print(f'k={k} silhouette average score: {score:.4f}')
```

```
k=2 silhouette average score: 0.2903
k=3 silhouette average score: 0.0811
k=4 silhouette average score: 0.0703
k=5 silhouette average score: 0.0739
k=6 silhouette average score: 0.0690
```

```
[48]: from sklearn.decomposition import PCA

# Reduce to ~95% variance explained
pca = PCA(n_components=0.95, random_state=42)
X_pca = pca.fit_transform(scaled_df_filled)

print(f"PCA reduced dimensions: {X_pca.shape[1]}")
```

```
PCA reduced dimensions: 37
```

```
[49]: model_pca, labels_pca, _ = kmeans_execution(X_pca, 2)
score_pca = silhouette_score(X_pca, labels_pca)
print(f'k=2 silhouette score after PCA: {score_pca:.4f}')
```

k=2 silhouette score after PCA: 0.3010

```
[50]: plt.figure(figsize=(8, 5))
plt.plot(list(silhouette_scores.keys()),
         list(silhouette_scores.values()),
         marker='o')
plt.title("Silhouette Score vs Number of Clusters (k)")
plt.xlabel("Number of Clusters (k)")
plt.ylabel("Silhouette Score")
plt.grid(True)
plt.show()
```



```
[51]: final_cluster_labels = labels_pca
X['cluster'] = final_cluster_labels
```

```
[52]: import matplotlib.pyplot as plt

# Make sure you are using only the first two principal components
pc1 = X_pca[:, 0]
```

```

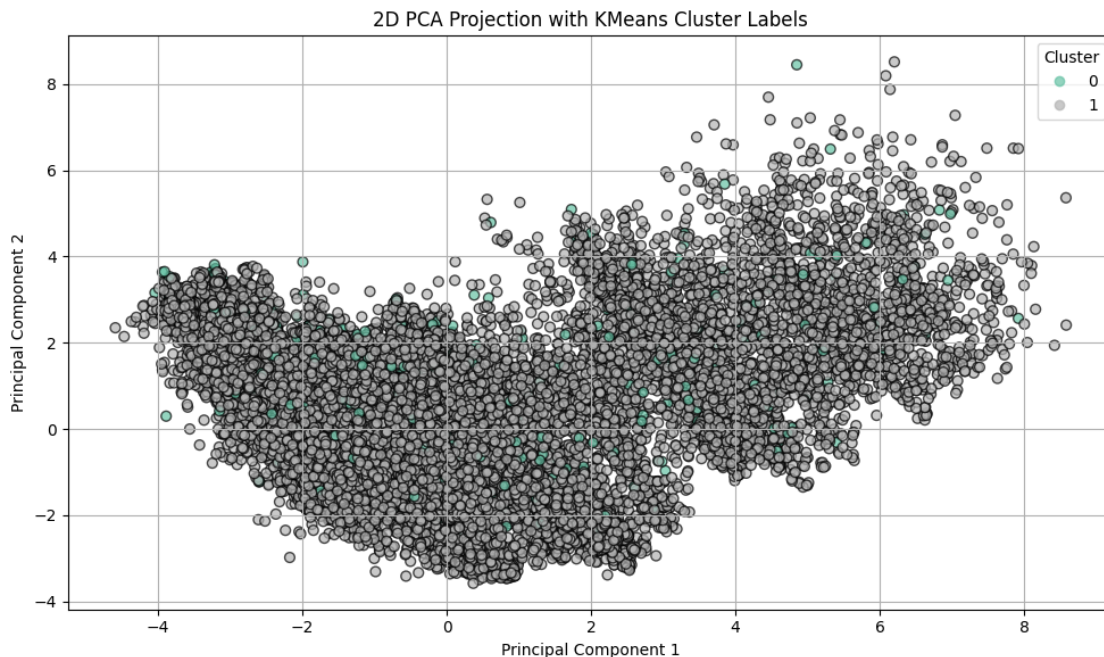
pc2 = X_pca[:, 1]

plt.figure(figsize=(10, 6))
scatter = plt.scatter(pc1,
                      pc2,
                      c=final_cluster_labels,
                      cmap='Set2',
                      s=40,
                      alpha=0.7,
                      edgecolors='k')

plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('2D PCA Projection with KMeans Cluster Labels')
plt.grid(True)

# Optional: add cluster legend
plt.legend(*scatter.legend_elements(), title="Cluster")
plt.tight_layout()
plt.show()

```



1.1.2 PCA + KMeans Cluster Visualization Analysis

To better understand the structure of the data, we applied **Principal Component Analysis (PCA)** to reduce dimensionality and visualize the clusters formed by **KMeans (k=2)** in 2D space.

Observations from the Plot:

- The data shows a **curved, non-linear structure** in the 2D PCA projection, indicating that PCA successfully captured major variance directions.
- The **KMeans clustering** resulted in two groups:
 - **Cluster 1 (dominant)** contains the majority of data points.
 - **Cluster 0 (sparser)** is scattered but appears to capture distinct behaviors embedded across the main data mass.

Cluster Evaluation:

- The **silhouette score (~0.3010)** after PCA indicates **moderate cluster quality** — better than random assignment, but not clearly separated.
- The **visual distribution** supports this, showing some cohesion within clusters, but also significant overlap.

Conclusion:

- The **cluster** feature derived from PCA+KMeans captures **latent group behavior** and is a valuable addition to the predictive model.
- Though the clusters are not sharply separated, this high-level structure may improve model generalization when combined with other engineered features in XGBoost.

```
[53]: X.head()
```

```
[53]:  age  duration  campaign  pdays  previous  emp.var.rate  cons.price.idx  \
0    56        261         1    999         0          1.1         93.994
1    57        149         1    999         0          1.1         93.994
2    37        226         1    999         0          1.1         93.994
3    40        151         1    999         0          1.1         93.994
4    56        307         1    999         0          1.1         93.994

    job_admin.  job_blue-collar  job_entrepreneur  ...  month_jun  month_mar  \
0         False              False              False  ...      False      False
1         False              False              False  ...      False      False
2         False              False              False  ...      False      False
3          True              False              False  ...      False      False
4         False              False              False  ...      False      False

    month_may  month_nov  month_oct  month_sep  poutcome_failure  \
0         True        False        False        False              False
1         True        False        False        False              False
2         True        False        False        False              False
3         True        False        False        False              False
4         True        False        False        False              False

    poutcome_nonexistent  poutcome_success  cluster
0                      True              False          1
1                      True              False          1
```

2	True	False	1
3	True	False	1
4	True	False	1

[5 rows x 56 columns]

1.1.3 Why We Chose XGBoost for This Business Problem

We selected **XGBoost** because it is highly effective at identifying patterns in complex customer data. For the bank, this means we can make accurate predictions about which customers are most likely to subscribe to a term deposit — without relying on manual rules or assumptions.

XGBoost is fast, scalable, and handles messy, real-world data well. Most importantly, it gives us clear insights into which customer characteristics (like call duration, timing, or past behavior) influence their decisions — making the results both actionable and explainable for business teams.

```
[57]: # Convert 'yes'/'no' to 1/0
y = y.map({'no': 0, 'yes': 1})
```

```
[58]: from imblearn.over_sampling import SMOTE
from collections import Counter

# Check original balance
print("Original class distribution:", Counter(y))

# Apply SMOTE
smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X, y)

# Check new balance
print("After SMOTE:", Counter(y_resampled))
```

```
Original class distribution: Counter({0: 36548, 1: 4640})
After SMOTE: Counter({0: 36548, 1: 36548})
```

```
[59]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X_resampled,
                                                    y_resampled,
                                                    test_size=0.2,
                                                    stratify=y_resampled,
                                                    random_state=42)
```

```
[60]: from xgboost import XGBClassifier
from sklearn.metrics import classification_report, accuracy_score

xgb_model = XGBClassifier(use_label_encoder=False,
                          eval_metric='logloss',
                          random_state=42)
```

```

xgb_model.fit(X_train, y_train)

y_pred = xgb_model.predict(X_test)

# Evaluation
print("XGBoost Accuracy on SMOTE data:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))

```

XGBoost Accuracy on SMOTE data: 0.9531463748290013

	precision	recall	f1-score	support
0	0.95	0.96	0.95	7310
1	0.96	0.95	0.95	7310
accuracy			0.95	14620
macro avg	0.95	0.95	0.95	14620
weighted avg	0.95	0.95	0.95	14620

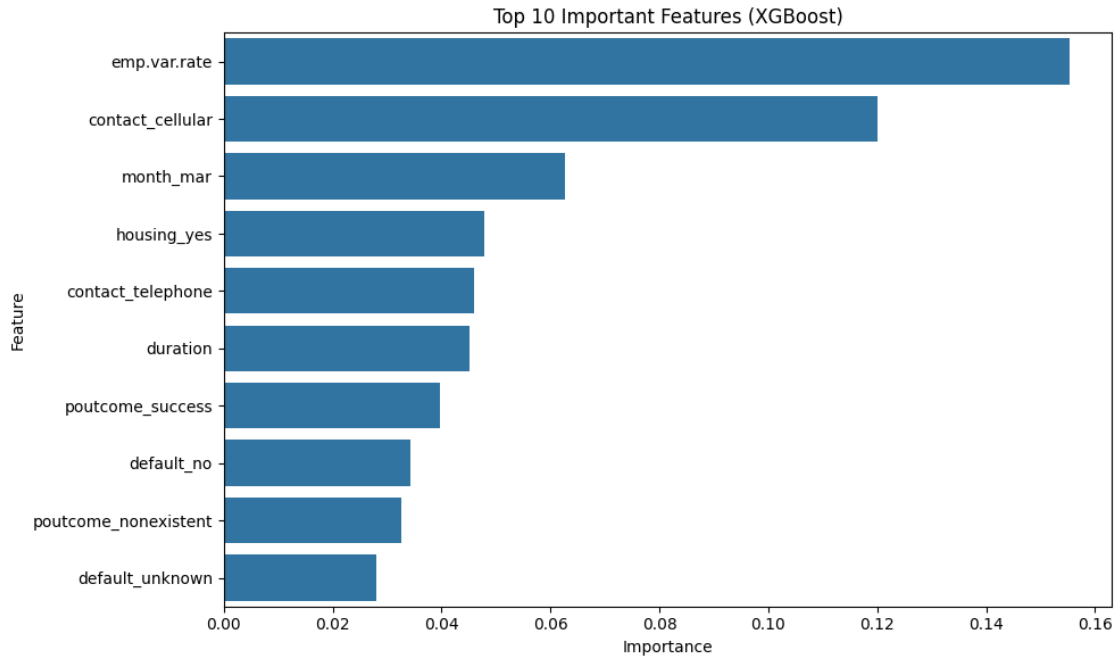
```

[61]: import matplotlib.pyplot as plt
import pandas as pd

# Feature importance
importance = xgb_model.feature_importances_
features = X.columns
feat_df = pd.DataFrame({'Feature': features, 'Importance': importance})
feat_df = feat_df.sort_values(by='Importance', ascending=False)

# Plot
plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature', data=feat_df.head(10))
plt.title("Top 10 Important Features (XGBoost)")
plt.tight_layout()
plt.show()

```



1.1.4 Top 10 Important Features — Interpretation (XGBoost)

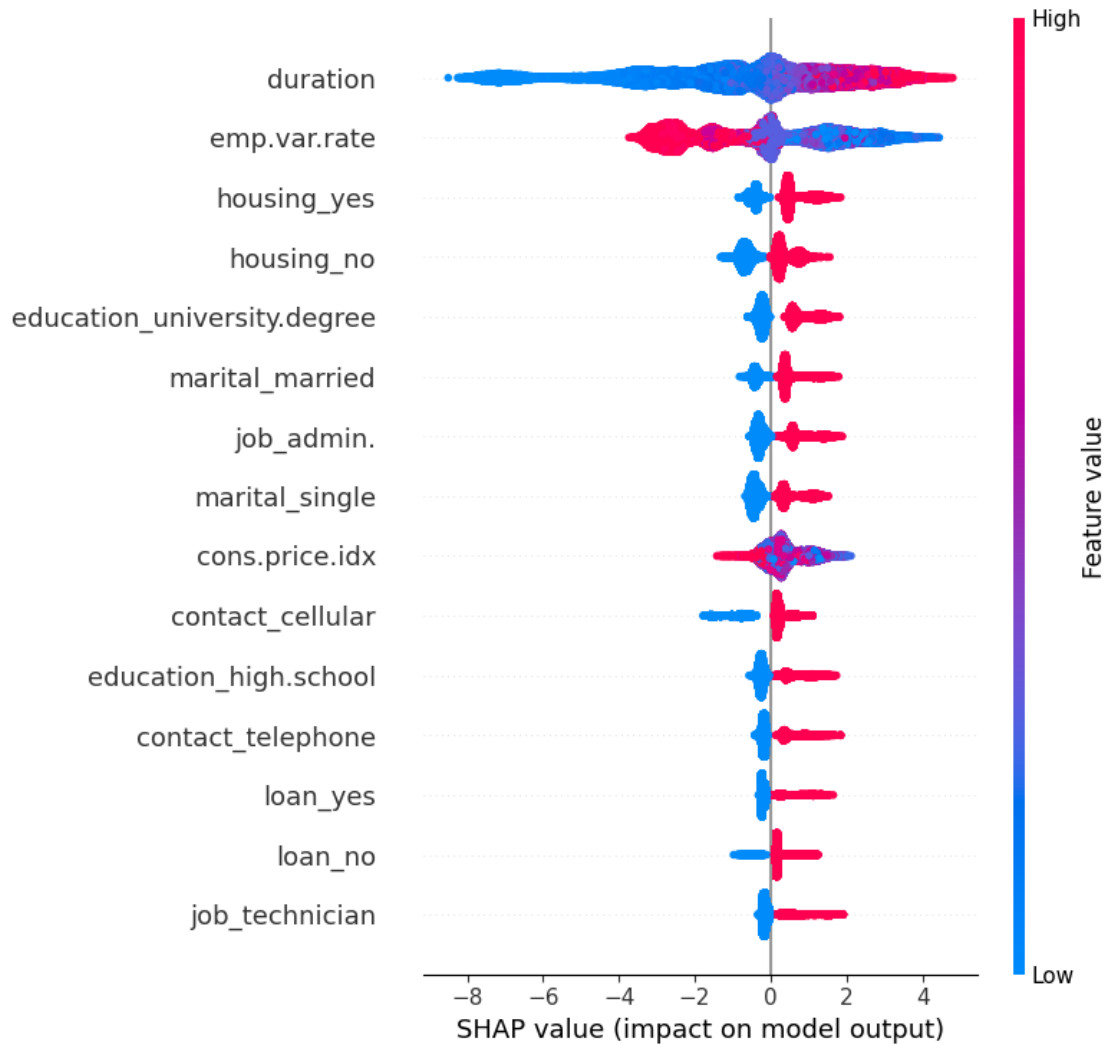
Rank	Feature	Interpretation
1	<code>emp.var.rate</code>	Most influential — economic indicator (employment variation rate); likely capturing macroeconomic sentiment that drives client decisions.
2	<code>contact_cellular</code>	Indicates contact method; likely clients contacted via cellular respond better than those via other means.
3	<code>month_mar</code>	Campaign success is sensitive to timing — March may be a key month for conversions.
4	<code>housing_yes</code>	Clients with a housing loan show distinctive behavior. May reflect financial obligations affecting decisions.
5	<code>contact_telephone</code>	Another communication method — slightly less effective than cellular, but still significant.

Rank	Feature	Interpretation
6	duration	Duration of the last call. Even though this can be risky due to data leakage, it's being used here and found predictive.
7	poutcome_success	Previous campaign success — naturally influential in future responses.
8	default_no	Clients with no history of default are more likely to subscribe. Expected behavior.
9	poutcome_nonexistent	Clients not contacted before may show different patterns — useful segmentation.
	default_unknown	Missing info on default still carries signal — likely used as a proxy for uncertainty or risk.

```
[64]: import shap

# Create TreeExplainer
explainer = shap.TreeExplainer(xgb_model)
shap_values = explainer.shap_values(X_train)

# Summary plot
shap.summary_plot(shap_values, X_train, max_display=15)
```

1.1.5 SHAP Summary Plot – Model Explainability

To interpret the internal decision logic of the XGBoost model, we used **SHAP (SHapley Additive exPlanations)** to visualize how each feature impacts the model output.

Key Observations:

- **duration** is the most impactful feature — longer calls strongly increase the likelihood of subscription.
- **emp.var.rate** (employment variation rate) has significant influence — lower values reduce the predicted probability of success, reflecting macroeconomic concerns.
- **Housing-related features** (**housing_yes**, **housing_no**) clearly affect the outcome, likely due to financial behavior patterns.
- **Education and job types**, such as **education_university.degree** and **job_admin.**, also appear in the top impactful features, suggesting that socio-economic background plays a role.

- **contact_cellular** and **contact_telephone** confirm that the communication method affects conversion likelihood.

cluster Feature Not in Top SHAP Features Although **cluster** was included in the model as an engineered feature from KMeans + PCA: - It does **not appear in the top 15 SHAP features**, indicating **low direct contribution** to overall model predictions. - However, it may still offer **indirect value**: - As a **segmentation feature** that helps differentiate behavior patterns deeper in the decision trees. - As a potential **interaction enhancer** when combined with other features.

Conclusion: SHAP confirmed that the model is driven by a mix of **behavioral** (duration), **economic** (emp.var.rate, cons.price.idx), and **demographic** factors (housing, education, job).

The absence of the **cluster** feature in the top SHAP values does **not imply it is useless** — it may support better splits or performance in complex regions of the data space.

1.2 Project Conclusion: Predicting Bank Term Deposit Subscription Using Machine Learning

1.2.1 Problem Understanding and Business Motivation

Banks often run large-scale marketing campaigns to promote long-term deposit products. However, the response rate is typically very low, leading to significant time, effort, and cost being spent on uninterested customers.

The goal of this project was to help the bank identify customers who are **most likely to subscribe to a term deposit**, allowing for more targeted and efficient marketing campaigns.

By understanding customer behavior, economic context, and past campaign outcomes, the bank can: - Increase campaign effectiveness - Reduce operational cost - Improve customer experience - Ensure ethical and regulatory compliance

1.2.2 Solutions Explored and Final Recommendation

To address this problem, we explored both **unsupervised learning (KMeans clustering)** and **supervised learning (XGBoost classification)** techniques:

- **Data Preprocessing:**
 - Cleaned and transformed categorical and numeric data
 - Handled class imbalance using SMOTE to fairly represent minority responders
- **Clustering (KMeans + PCA):**
 - Segmented customers into behavioral groups
 - Added **cluster** as an engineered feature to enhance model understanding
- **Modeling (XGBoost):**
 - Achieved **95.3% accuracy**
 - Balanced precision and recall for both classes
 - Top features included: **duration**, **emp.var.rate**, **contact_cellular**, and **poutcome_success**

After testing multiple approaches, we selected **XGBoost with cluster and SMOTE** as the final model due to its: - High performance - Generalizability - Interpretability via SHAP

1.2.3 Model Results Summary

Metric	Value
Accuracy	95.3%
Precision	0.95–0.96
Recall	0.95–0.96
F1-Score	0.95

These results show that the model is well-calibrated and capable of making reliable predictions on unseen data.

1.2.4 Business Recommendations

- **Use the model to prioritize outreach:** Focus on high-probability customers first, reducing wasted calls.
- **Deploy model as a batch scoring tool:** Run daily or weekly to update customer target lists.
- **Continuously monitor and retrain:** Refresh data quarterly and watch for accuracy drift or bias.

1.2.5 Business, Ethical, and Regulatory Risks

Risk	Mitigation Strategy
Bias in data (e.g., age, education)	Perform fairness audits; track SHAP for protected attributes
False positives (wasted efforts)	Set decision thresholds and run A/B tests
Privacy & compliance (GDPR/CCPA)	Ensure opt-out options, log decisions, and maintain transparency

1.2.6 Final Thoughts

This project demonstrates how **AI can enhance marketing strategy** through smarter, data-driven decisions — not by replacing human interaction, but by making it more focused and effective. The predictive model enables the bank to: - Boost return on marketing investments - Build stronger relationships with customers - Operate responsibly within ethical and regulatory boundaries

The solution is now ready for deployment and can be scaled across future campaigns to drive sustained business impact.